

R. Thomas McCoy

Department of Cognitive Science
141 Krieger Hall
Johns Hopkins University
3400 N. Charles Street
Baltimore, MD 21218-2685

Email: tom.mccoy@jhu.edu
Last updated: August 20, 2021

RESEARCH INTERESTS

I study computational linguistics using techniques from cognitive science, machine learning, and natural language processing. My research focuses on how to achieve robust generalization in models of language, as this remains one of the main areas where current AI systems fall short and one of the most impressive components of language processing in humans. In particular, I study which **inductive biases** and which **representations of structure** enable robust generalization, since these are two of the major components that determine how learners generalize to novel types of input. Underlying much of this work is an additional theme of **analyzing and interpreting** neural network models.

EDUCATION

- | | |
|--------------|--|
| 2017–present | Johns Hopkins University: Ph.D. in Cognitive Science. GPA: 4.0.
<i>Advisors:</i> Tal Linzen, Paul Smolensky |
| 2013–2017 | Yale University: B.A. in Linguistics, <i>summa cum laude</i> , distinction in the major. GPA: 4.0.
<i>Advisor:</i> Robert Frank |
| Summer 2016 | Institute on Collaborative Language Research (CoLang), University of Alaska Fairbanks |
| Summer 2015 | Linguistic Summer Institute, University of Chicago |

EMPLOYMENT

- | | |
|-------------|--|
| Summer 2020 | Microsoft Research intern
Supervisor: Asli Celikyilmaz
<i>Evaluation methods for neural text generation systems</i> |
| Summer 2018 | Jelinek Summer Workshop on Speech and Language Technology (JSALT) sentence representations team
Team leaders: Sam Bowman, Ellie Pavlick
<i>Analysis techniques for learned sentence representations.</i> |

- Summer 2017 Carnegie Mellon University team for the DARPA Low Resource Languages for Emergent Incidents (LORELEI) project
Team leader: Patrick Littell
Finite-state morphological analyzer for Oromo.
- Summer 2017 Chirila project
Supervisor: Claire Bower
Developed automatic semantic processing techniques for an online database of Australian languages.
- Summer 2016 Grammar Boot Camp
Supervisor: Claire Bower
Wrote a sketch grammar of Kuwarra.
- Summer 2015 Yale Grammatical Diversity Project
Supervisors: Laurence Horn, Jim Wood, Raffaella Zanuttini, Jason Zentz
Edited web pages about regional grammatical phenomena.
- Summer 2014 Irish lip rounding research
Supervisor: Ryan Bennett
Collected lip rounding measurements from images of Irish speakers.
- Summer 2014 Linguistic Core Multi-University Research Initiative
Team leaders: Chris Dyer, Lori Levin
English-to-Malagasy tree-to-string transducer.
- Summer 2013 Linguistic Core Multi-University Research Initiative
Team leaders: Chris Dyer, Lori Levin
Finite state morphological analyzer for Kinyarwanda.

TEACHING

- Spring 2020 Johns Hopkins University
Role: Teaching Assistant
Course: Foundations of Cognitive Science
Lecture Instructor: Paul Smolensky
Led one seminar discussion and graded assignments.
- Fall 2019 Johns Hopkins University
Role: Teaching Assistant, Lab Instructor
Course: Computational Psycholinguistics
Lecture Instructor: Tal Linzen
Led lab sessions and graded assignments.
- Spring 2019 Johns Hopkins University
Role: Teaching Assistant
Course: Syntax I
Lecture Instructor: Géraldine Legendre
Led review sessions and graded assignments.

- Fall 2018 Johns Hopkins University
 Role: Teaching Assistant
 Course: Introduction to Computational Cognitive Science
 Lecture Instructor: Tal Linzen
Created educational simulations, tutorials, and homeworks in Javascript and Jupyter and taught lectures using these resources.
- Spring 2018 Johns Hopkins University
 Role: Fieldwork Instructor
 Course: World of Language
 Lecture Instructor: Géraldine Legendre
Led two sections of weekly fieldwork sessions complementing lectures.
- Summer 2015 Linguistic Society of America Summer Institute
 Role: Workshop Co-Instructor
 Course: Linguistic Enigmatography
 Co-Instructor: Lori Levin
Developed and co-taught a one-week workshop on creating linguistic puzzles.

AWARDS

1. Fellowships

- 2018–2021 NSF Graduate Research Fellowship
Project title: Assessing the capacity of computational models to make linguistic generalizations
- 2021 Sweitzer Fellow
Fellowship awarded by the Johns Hopkins Department of Cognitive Science to one graduate student.
- 2020 Finalist: Facebook Fellowship
One of four finalists in the Natural Language Processing category; two of the four finalists received fellowships.
- 2017–2020 Owen Scholars Fellowship
Fellowship for outstanding incoming Johns Hopkins PhD students in the natural sciences.
- 2017 Finalist, Rhodes Scholarship
- 2017 Finalist, Marshall Scholarship

2. Prizes

- 2017 Alpheus Henry Snow Prize
Award for the graduating Yale senior who is “adjudged by the faculty to have done the most for Yale by inspiring in his or her classmates an admiration and love for the best traditions of high scholarship.”

- 2016 Hart Lyman Prize
Award for the Yale junior who “has made through his/her own efforts the best record intellectually and socially.”
- 2016 Phi Beta Kappa
One of 13 Yale students admitted as juniors.
- 2013 World champion team at the International Linguistics Olympiad
Member of the four-person U.S. team selected by the North American Computational Linguistics Olympiad.
- 2013 United States Presidential Scholar
One of two for Pennsylvania.

3. Grants

- 2019 NeurIPS Travel Grant
Grant to fund travel to present work at the NeurIPS workshop on Context and Compositionality in Biological and Artificial Neural Systems.
- 2019 ICLR Travel Grant
Grant to fund travel to present two projects at the 2019 ICLR conference.
- 2018–2019 Johns Hopkins University Center for Educational Resources Technology Fellowship Grant
Co-Grantee: Tal Linzen
Grant to develop interactive visualizations of concepts in computational cognitive science.

MAIN PROJECTS

See the PUBLICATIONS section below for the full references of the papers referred to here.

2018–present: Discovering the compositional structure implicitly learned by neural networks

- Implemented several novel architectures based in the Tensor Product Representation formalism as a tool for analyzing the vector representations learned by neural network models. The goal of this analysis is to test the hypothesis that these vector representations are implicitly symbolic structures. GitHub: <https://github.com/tommccoy1/tpdn>
- Analyzed the representations of models trained on synthetic string-manipulation tasks as well as state-of-the-art sentence embedding models. Published in [10] and [8]. Further results are being written up in [26].
- Created a Javascript demo: https://rtmccoy.com/tpdn/tpr_demo.html
- This project is the topic of my dissertation, which is ongoing.

2020–present: Evaluating novelty in neural text generation systems

- Internship project at Microsoft Research analyzing the extent to which text generated by GPT-2, a state-of-the-art text generation system, is novel vs. copied from the training set.
- Results and code are ready for release in [24].

2020–present: Imparting linguistic inductive biases via meta-learning

- Implemented model-agnostic meta-learning (MAML) in PyTorch and applied it to neural networks to give them learning biases based on Optimality Theory. GitHub: <https://github.com/tommccoy1/meta-learning-linguistic-biases>.
- Analyzed data efficiency and generalization for a model that had undergone meta-learning. Published the results in [2].
- JavaScript demo: <http://rtmccoy.com/meta-learning-linguistic-biases.html>

2019–present: Artificial language learning experiments testing people’s syntactic generalization

- Conducted an Artificial Language Learning experiment on Mechanical Turk testing how people generalize center embedding. Published the results in [1].
- In ongoing work, we have expanded the Bayesian ideal learner model of Perfors, Tenenbaum, and Regier (2011) by enabling it to handle unification constraints. We are now using this augmented model to design additional experiments.

2017–2020: Understanding syntactic inductive biases of sequence-to-sequence recurrent neural networks

- Implemented many types of sequential and tree-based recurrent neural networks in PyTorch, including novel variations of these architectures. GitHub: <https://github.com/tommccoy1/rnn-biases>
- Trained models on syntactic tasks such as question formation (e.g., changing *The man is happy* into *Is the man happy?*), and analyzed how models’ syntactic inductive biases are affected by model structure, hyperparameters, and training setup (e.g., single-task vs. multi-task learning).
- Published results in [3] and [6]; further results are on the project website: http://rtmccoy.com/rnn_hierarchical_biases.html.
- In work that is in preparation, we analyzed models trained on the CHILDES corpus to test which learning biases are sufficient to induce models to generalize as children do from the type of input that children receive, to be submitted as [25].

2019–present: Heuristic Analysis of NLI Systems

- Created a controlled evaluation set called Heuristic Analysis of NLI Systems (HANS) which tests how robustly models trained on natural language inference (NLI) generalize to examples outside of their training distribution. GitHub: <https://github.com/tommccoy1/hans>
- Evaluated 4 state-of-the-art models on the HANS dataset, revealing how even these models are overly reliant on heuristics rather than valid inference. Published in [16].
- Conducted follow-up work showing how different random restarts of BERT can have vastly differing performance on HANS. Published in [5].
- In ongoing work, we are studying data augmentation as a solution to these problems. Some of this work was published in [15].

2018: JSALT sentence representations team

- Member of a team at the Jelinek Summer Workshop on Speech and Language Technology (JSALT) studying learned sentence representations.
- Evaluated various state-of-the-art models on syntax-based tasks, and analyzed their performance on these tasks. Resulted in components of [9], [11], and [17].

Inductive biases in humans and machines

- [1] 2021 **R. Thomas McCoy**, Jennifer Culbertson, Paul Smolensky, and G  rardine Legendre. [Infinite use of finite means? Evaluating the generalization of center embedding learned from an artificial grammar](#). In *CogSci Conference 2021*. Poster.
- [2] 2020 **R. Thomas McCoy**, Erin Grant, Paul Smolensky, Thomas L. Griffiths, and Tal Linzen. [Universal linguistic inductive biases via meta-learning](#). In *CogSci Conference 2020*. Talk.
- [3] 2020 **R. Thomas McCoy**, Robert Frank, and Tal Linzen. [Does syntax need to grow on trees? Sources of hierarchical inductive bias in sequence-to-sequence networks](#). *TACL*. Journal article.
- [4] 2020 Michael Lepori, Tal Linzen, and **R. Thomas McCoy**. [Representations of Syntax \[MASK\] Useful: Effects of Constituency and Dependency Structure in Recursive LSTMs](#). *ACL*. Pre-recorded virtual talk.
- [5] 2020 **R. Thomas McCoy**, Junghyun Min, and Tal Linzen. [BERTs of a feather do not generalize together: Large variability in generalization across models with similar test set performance](#). *BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*.
- [6] 2018 **R. Thomas McCoy**, Robert Frank, and Tal Linzen. [Revisiting the poverty of the stimulus: hierarchical generalization without a hierarchical bias in recurrent neural networks](#). In *CogSci Conference 2018*.

Representations of structure

- [7] 2020 Michael Lepori and **R. Thomas McCoy**. [Picking BERT’s brain: Analyzing contextualized embeddings using Representational Similarity Analysis](#). *COLING*.
- [8] 2020 Paul Soulos, **R. Thomas McCoy**, Tal Linzen, and Paul Smolensky. [Uncovering the compositional structure of vector representations with Role Learning Networks](#). *BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*.
- [9] 2019 Najoung Kim, Roma Patel, Adam Poliak, Alex Wang, Patrick Xia, **R. Thomas McCoy**, Ian Tenney, Alexis Ross, Tal Linzen, Benjamin Van Durme, Samuel R. Bowman, Ellie Pavlick. [Probing What Different NLP Tasks Teach Machines about Function Word Comprehension](#). *Proceedings of the Eighth Joint Conference on Lexical and Computational Semantics (*SEM 2019)*. Talk.
Best paper award at *SEM 2019.
- [10] 2019 **R. Thomas McCoy**, Tal Linzen, Ewan Dunbar, and Paul Smolensky. [RNNs implicitly implement tensor-product representations](#). *ICLR 2019*.

- [11] 2019 Ian Tenney, Patrick Xia, Berlin Chen, Alex Wang, Adam Poliak, **R. Thomas McCoy**, Najoung Kim, Benjamin Van Durme, Samuel R. Bowman, Dipanjan Das, and Ellie Pavlick. [What do you learn from context? Probing for sentence structure in contextualized word representations.](#) *ICLR 2019*.
- [12] 2017 Jungo Kasai, Robert Frank, **R. Thomas McCoy**, Owen Rambow, and Alexis Nasr. [TAG parsing with neural networks and vector representations of supertags.](#) *EMNLP 2017*.
- [13] 2017 Dan Friedman*, Jungo Kasai*, **R. Thomas McCoy***, Robert Frank, Forrest Davis, and Owen Rambow. [Linguistically Rich Vector Representations of Supertags for TAG Parsing.](#) In *Proceedings of the 13th International Workshop on Tree Adjoining Grammars and Related Formalisms*. Talk.
*Equal contribution.

Robust generalization

- [14] 2021 Paul Soulos, Sudha Rao, Caitlin Smith, Eric Rosen, Asli Celikyilmaz, **R. Thomas McCoy**, Yichen Jiang, Coleman Haley, Roland Fernandez, Hamid Palangi, Jianfeng Gao and Paul Smolensky. [Structural Biases for Improving Transformers on Translation into Morphologically Rich Languages.](#) *LoResMT 2021*. Pre-recorded virtual talk.
- [15] 2019 Junghyun Min, **R. Thomas McCoy**, Dipanjan Das, Emily Pitler, and Tal Linzen. [Syntactic data augmentation increases robustness to inference heuristics.](#) *ACL 2020*. Pre-recorded virtual talk.
- [16] 2019 **R. Thomas McCoy**, Ellie Pavlick, and Tal Linzen. [Right for the Wrong Reasons: Diagnosing Syntactic Heuristics in Natural Language Inference.](#) *ACL 2019*. Talk.
- [17] 2019 Samuel R. Bowman, Ellie Pavlick, Edouard Grave, Benjamin Van Durme, Alex Wang, Jan Hula, Patrick Xia, Raghavendra Pappagari, **R. Thomas McCoy**, Roma Patel, Najoung Kim, Ian Tenney, Yinghui Huang, Katherin Yu, Shuning Jin, and Berlin Chen. [Can You Tell Me How to Get Past Sesame Street? Sentence-Level Pretraining Beyond Language Modeling.](#) *ACL 2019*.
- [18] 2019 **R. Thomas McCoy** and Tal Linzen. [Non-entailed subsequences as a challenge for natural language inference.](#) In *Proceedings of the Society for Computation in Linguistics (SCiL) 2019*.

Other topics

- [19] 2019 **R. Thomas McCoy**. [Touch down in Pittsburghese.](#) *Yale Working Papers in Grammatical Diversity*. Working paper.

- [20] 2018 Patrick Littell, **R. Thomas McCoy**, Na-Rae Han, Shruti Rijhwani, Zaid Sheikh, David Mortensen, Teruko Mitamura, and Lori Levin. [Parser combinators for Tigrinya and Oromo morphology](#). In *Language Resources and Evaluation Conference (LREC) 2018*.
- [21] 2018 **R. Thomas McCoy** and Robert Frank. [Phonologically Informed Edit Distance Algorithms for Word Alignment with Low-Resource Languages](#). In *Proceedings of the Society for Computation in Linguistics (SCiL) 2018*.
- [22] 2017 **R. Thomas McCoy**. [English comparatives as degree-phrase relative clauses](#). In *Proceedings of the Linguistic Society of America 2*.

WORK IN PREPARATION

- [23] Paul Smolensky, **R. Thomas McCoy**, Roland Fernandez, Matthew Goldrick, and Jianfeng Gao. The new generation of AI systems empowered by key cognitive principles. Accepted to *AI Magazine*.
- [24] **R. Thomas McCoy**, Paul Smolensky, Tal Linzen, Jianfeng Gao, and Asli Celikyilmaz. How much do language models copy from their training data? Evaluating linguistic novelty in text generation. In preparation to submit to *TACL*.
- [25] Aditya Yedetore, Tal Linzen, Robert Frank, and **R. Thomas McCoy**. How poor is the stimulus? Evaluating hierarchical generalization in neural networks trained on child-directed speech. In preparation to submit to *ACL*.
- [26] **R. Thomas McCoy**, Tal Linzen, and Paul Smolensky. DISCOVER: A framework for dissecting compositionality in vector representations. In preparation to submit to *JAIR*.

UNPUBLISHED CONFERENCE PRESENTATIONS

- 2018 R. Thomas McCoy, Robert Frank, and Tal Linzen. Investigating hierarchical bias in the acquisition of English question formation with recurrent neural networks. Poster presentation, *2018 Legrain conference: Learning Language in Humans and in Machines*, Paris, France, July 5-6.
- 2018 Robert Frank, R. Thomas McCoy, and Tal Linzen. Neural network syntax in the age of deep learning: the case of question formation. Oral presentation, *Society for Computation in Linguistics*, Salt Lake City, Utah, January 5.
- 2017 Patrick Littell, R. Thomas McCoy, and Lori Levin. The North American Computational Linguistics Olympiad. Oral presentation, in *Datablitz: Getting High School Students into Linguistics: Current Activities and Future Directions*, *Linguistic Society of America Annual Meeting*, Austin, Texas, January 7.

INVITED TALKS

- 2021 Montreal Computational and Quantitative Linguistics Lab (MCQLL). October 26, 2021.
Discovering implicit compositional representations in neural networks
- 2021 Edinburgh Centre for Language Evolution. September 28, 2021.
How do neural networks represent compositional symbolic structure?
- 2021 USC ISI Natural Language Seminar. April 15, 2021.
Universal Linguistic Inductive Biases via Meta-Learning.
- 2020 DeepMind language reading group. December 7, 2020.
Analyzing the syntactic inductive biases of sequence-to-sequence networks.
- 2020 Berkeley NLP Seminar. October 16, 2020.
Analyzing the syntactic inductive biases of sequence-to-sequence networks.
- 2020 NLP With Friends seminar series. August 12, 2020.
Universal Linguistic Inductive Biases via Meta-Learning.
- 2019 Workshop on Gradient Symbolic Computation. Johns Hopkins University. September 19, 2019.
Tensor product decomposition of continuous vector representations
- 2018 Microsoft Research, Redmond. December 11, 2018.
Discovering the compositional structure implicitly learned by neural networks

MENTORING

Master's students

- 2019–2020 Junghyun Min
Co-authored [5] and [15]. Co-supervised with Tal Linzen.
- 2019–2020 Paul Soulos
Co-authored [8]. Co-supervised with Paul Smolensky.

Undergraduate students

- 2019–present Aditya Yedetore
Ongoing research to be submitted as [25]. Co-supervised with Tal Linzen.
- 2019–2020 Michael Lepori
Co-authored [4] and [7]. Co-supervised with Tal Linzen.

SERVICE

- 2020–present Johns Hopkins CogSci Representation and Diversity Committee. Co-created and co-organized a program for giving prospective students feedback on their PhD applications.
- 2016–2017 Computational Linguistics at Yale (CLAY) reading group: Co-organizer.
- 2015–2017 Yale Undergraduate Linguistics Society: Co-founder (2015), president (2015–2016), treasurer (2016–2017).

REVIEWING

- 2021 Workshop reviewer: BlackboxNLP 2021.
- 2021 Journal reviewer: Natural Language Engineering.
- 2021 Conference reviewer: EMNLP 2021.
- 2020 Conference reviewer: CoNLL 2020.
- 2020 Conference reviewer: EMNLP 2020. Recognized as an outstanding reviewer.
- 2020 Conference reviewer: ACL 2020.
- 2019 Conference reviewer: CoNLL 2019.
- 2018 Conference reviewer: CoNLL 2018.
- 2018 Conference reviewer: ACL 2018. Recognized as a top reviewer.

OUTREACH

- 2020 Public talk for the National Museum of Language: *Language Squared: The Linguistics of Crosswords*.
- 2013–present North American Computational Linguistics Olympiad (NACLO). National level: Co-Program Chair; problem writer (16 problems to date). Local level: Co-founder and co-organizer of the Yale contest site (2013–2017); co-organizer of the Johns Hopkins contest site (2017–present); organizer of pre-contest practice sessions at both sites.
- 2018–2019 International Linguistics Olympiad (IOL): Problem writer.
- 2016 Yale Grammatical Diversity Project: Authored two webpages describing regional grammatical phenomena (*All the further* and *Subject contact relatives*).
- 2013–2017 Linguistics teaching initiatives: Designed and taught a one-lecture linguistics class to high school students in connection with the separate programs Splash, Sprout, and Math Mornings. Presented 8 times to groups ranging from 25 to 50 students.

PROFESSIONAL MEMBERSHIPS

- 2015–present Linguistic Society of America (LSA).
- 2017–present Association for Computational Linguistics (ACL).
- 2018–present Cognitive Science Society.

SKILLS

Programming languages Python, PyTorch, JavaScript, Haskell, C, Java, R, Scheme.

Natural languages English (native), Bahasa Indonesia (basic conversation), Old English (basic reading ability), Old Norse (basic reading ability), Latin (basic reading ability).

COURSEWORK

Undergraduate GPA: 4.0 Graduate GPA: 4.0

Computational Linguistics: Language and Computation I, Language and Computation II, Formal Foundations of Linguistic Theories, Computing Meaning

Natural Language Processing: Natural Language Processing, Machine Learning: Linguistic and Sequence Modeling

Syntax: Syntax I, Syntax II, Grammatical Diversity in US English

Phonetics/Phonology: Phonetics, Phonology I, Phonology II, The Phonetics/Phonology Interface

Semantics: Semantics I, Semantics II

Computer Science: Data Structures and Programming Techniques, Computational Tools for Data Science

Mathematics: Multivariable Calculus, Discrete Mathematics, Probability and Statistics, Advanced Statistical Methods

Other relevant courses: Linguistic Field Methods, Foundations of Cognitive Science